



# Improving inference about cognitive processes using mixture models.

Gidon Frischkorn & Vencislav Popov





## Workshop Agenda

- I. What are Mixture Models and why can they be useful?
- II. Specifying mixture models in brms
  - a) Data formatting
  - b) Setting up mixture families
  - c) Understanding & Identifying parameters of mixture families
  - d) Fitting & Summarizing results of mixture models
- III. bmm – Easy implementation of mixture models for visual working memory tasks

--- Coffee Break ---

- IV. Work with (your own) data
- V. Outlook: Specifying custom mixture models for accuracy

## What will you (not) learn today?



- How to specify simple mixture models in brms
- How to use the *bmm* package to fit existing mixture models for visual working memory tasks
  - Two-parameter (Zhang & Luck, 2008)
  - Three-parameter (Bays et al., 2009)
  - Different flavors of the Interference Measurement model (Oberauer & Lin, 2017)
- Interpret & Summarize results of mixture models estimated using brms




- How to specify complicated mixture models or develop entirely new models
- How to fit mixture models for groups of subjects that differ in their behavior



## What are Mixture Models and why can they be useful?

**Assumptions** in (standard) data analysis:

1. DV stems from a single distribution (oftentimes a normal distribution)
2. predict parameters (usually the mean) from these distributions by independent variables




# What are Mixture Models and why can they be useful?

**Assumptions** in (standard) data analysis:

1. DV stems from a single distribution (oftentimes a normal distribution)
2. predict parameters (usually the mean) from these distributions by independent variables

## Problem:

→ Sometime data do stem from multiple different distributions






# What are Mixture Models and why can they be useful?

## Mixture Models...

- ...specify a set of distributions that data can stem from
- ...allow to estimate what proportion of data stems from each distribution
- ...enable to predict parameters of the different distributions

Such mixtures can occur...

- ...within a participant (different cognitive states or sources of signal)
- ...between participants



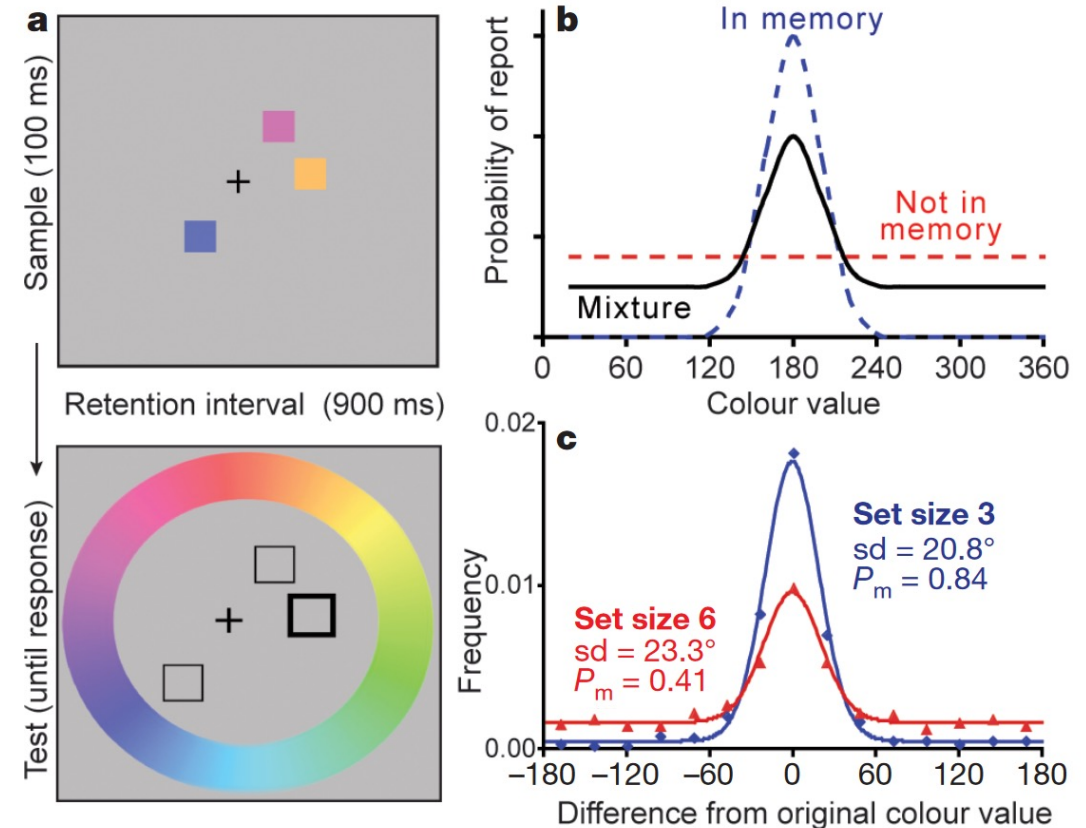
## What are Mixture Models and why can they be useful?

### Application to visual working memory

- Theories separate different states we can be at during retrieval
- a) encoded the item in memory → retrieval with the precision of memory representations
  - b) not encoded the item in memory → random guessing

Performance indicators:

1.  $P_{\text{mem}}$  = probability of having an item in memory
2. Precision = deviations from correct item



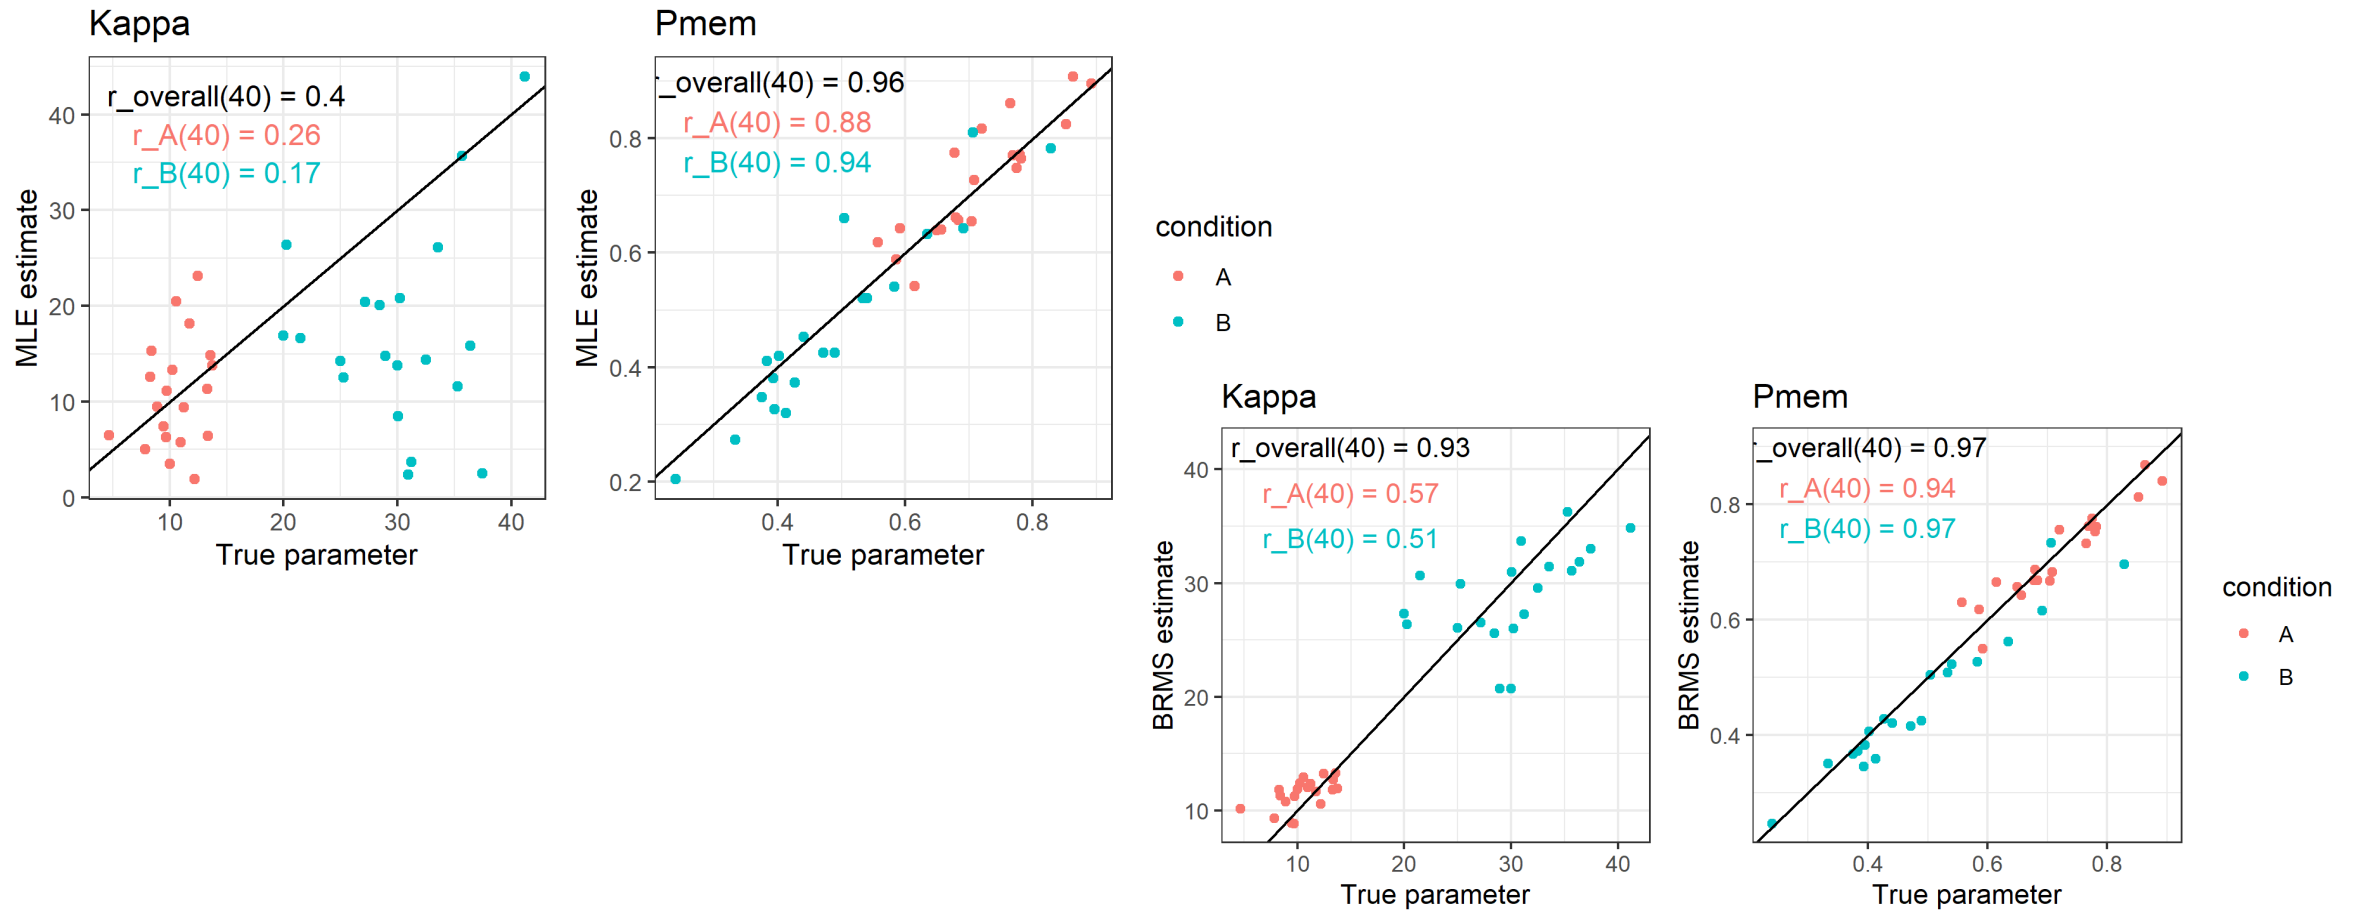



## brms/bmm in R vs MemToolbox in Matlab

	brms/bmm	MemToolbox
Estimation	Bayesian	Bayesian and Maximum Likelihood
Fitting multiple conditions	Jointly (Linear model syntax)	Separately to each condition
Inference over multiple conditions	1-step procedure	2-step procedure
Allows continuous predictors	Yes	No
Can fix some parameters across conditions	Yes	No
Tasks	Continuous report, custom	Continuous report, Change detection
Included models “out-of-the-box”	2-parameter, 3-parameter, Interference measurement model	2-parameter, 3-parameter, Variable precision, Slot+averaging, Slots+resources
Can customize models	Yes	No



## brms/bmm in R vs MemToolbox in Matlab






## Specifying mixture models in brms

### A (short) intro to brms (Bayesian regression models using Stan)

- I. interface to Stan → fit Bayesian generalized linear models
- II. formula syntax similar to lme4 → provides familiar & simple regression analyses.
- III. wide range of response distributions supported → fit wide range of data
- IV. lots of further modeling options:
  - non-linear and smooth terms,
  - auto-correlation structures,
  - censored data,
  - missing value imputation,
  - and quite a few more (like mixture models 😊)
- V. all parameters of response distribution can be predicted (means, standard deviations, etc.)
- VI. flexible prior specifications → encourages users to apply prior distributions that reflect their beliefs.
- VII. Model fit can easily be assessed → posterior predictive checks, cross-validation, and Bayes factors.



## Specifying mixture models in brms

### A (short) intro to brms (Bayesian regression models using Stan)

`brm` → main function to fit models using brms

#### required arguments:

1. **formula** → specifies the regression model to estimate
2. **data** → data set that contains all variables (**important:** match variable names to names in formula!)

#### defaults:

3. **family** = `gaussian()` → which probability distribution does the DV stem from
4. **priors**
5. **sampler settings** (number of warmup & post-warmup sample, number of MCMC chains, etc.)

```
brm(formula = DV ~ 1 + IV + (1 + IV | ID),  
    data = myData,  
    # optional arguments  
    family = gaussian(),  
    prior = myPriors,  
    iter = 4000, warmup = 1500,  
    nChains = 4)
```



## Specifying mixture models in brms

### Setting up mixture families

brms allows to specify mixtures of any of the supported data distributions 😊

Steps when setting up mixture models in *brms*:

1. Specify the mixture family
2. Specify model formula to predict parameters
3. Set priors to identify the different mixture components
4. Estimate the model (and have some patience!)
5. Evaluate results

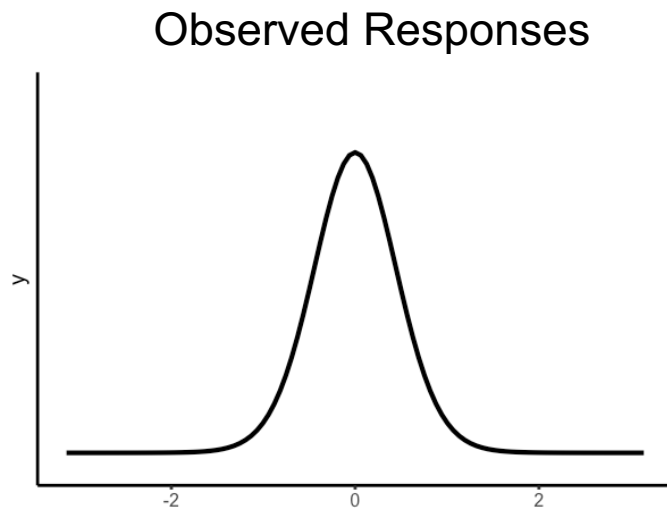
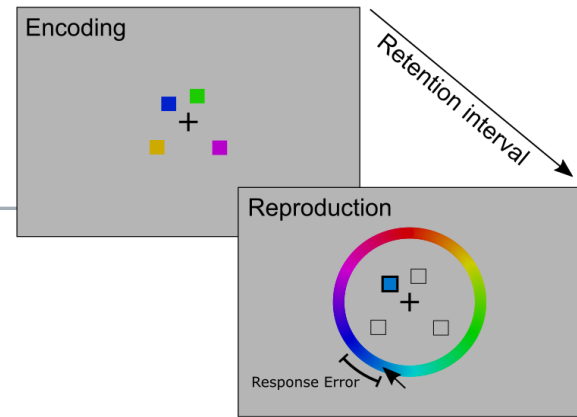
#### Option

If you want to follow along the next steps directly in R, open the R script „**2pMM\_Zhang&Luck2008\_brms.R**“.

This script implements all steps that we address now, using only brms syntax. Please note that there are some additional things implemented in the script, that are not included in the here.

# Specifying mixture models in brms

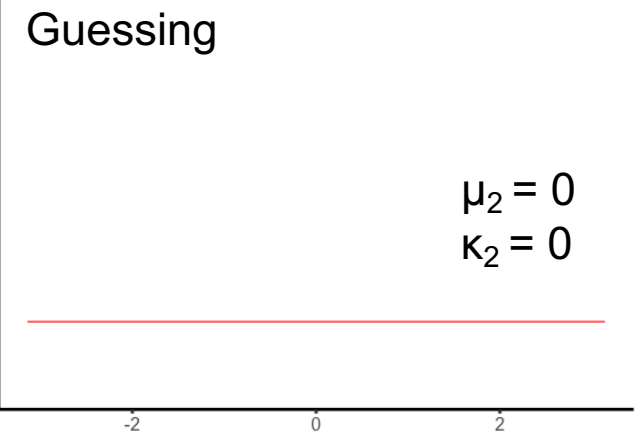

## Understanding & Identifying parameters of mixture families



$\theta_1 \rightarrow \text{estimated}$

$\theta_2 = 0$

$$\text{Softmax: } p_i = \frac{e^{\theta_i}}{\sum_{j=1}^K e^{\theta_j}}$$





## Specifying mixture models in brms

### Setting up mixture families

#### Step 1: Specifying mixture families in brms

→ `mixture()` function allows setting up mixtures of any set of distributions implemented in *brms*

Example: Two-Parameter Mixture model

```
mixFamily_ZL <- mixture(von_mises, von_mises)
```

#### Parameters for Mixture Families:

- parameters of each mixture distribution :
  - 1) vonMises<sub>1</sub>:  $\mu_1$  &  $\kappa_1$
  - 2) vonMises<sub>2</sub>:  $\mu_2$  &  $\kappa_2$
- mixing proportions
  - Each mixture distribution gets a mixing proportion →  $\theta$
  - mixing proportions are converted to probabilities → Softmax
  - One mixing proportion needs to be fixed for scaling

## Specifying mixture models in brms

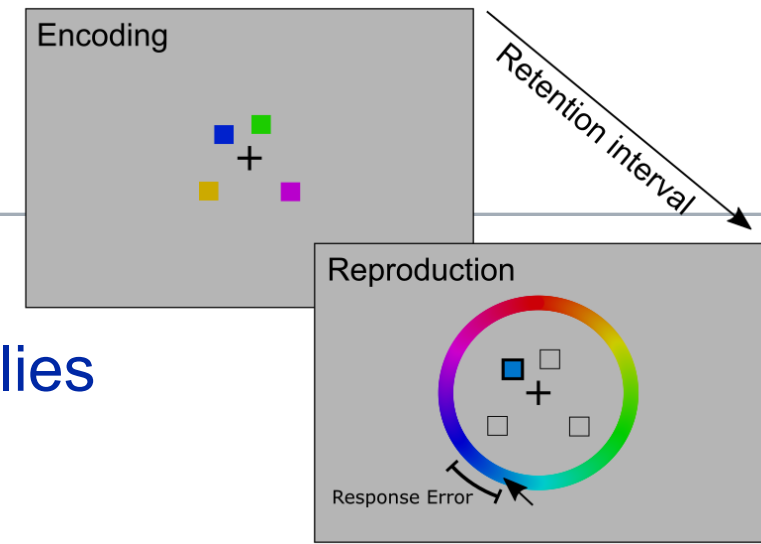
### Understanding & Identifying parameters of mixture families

#### Step 2: Specifying the model formula

```
formula_ZL <- brmsformula(ResponseError ~ 1,  
                             kappa1 ~ 1 + setsize + (1 + setsize | ID),  
                             mu2 ~ 1,  
                             kappa2 ~ 1,  
                             theta1 ~ 1 + setsize + (1 + setsize | ID))
```

#### Formula elements:

- Declaring the **dependent variable** → Response Error
- Predicting **estimated parameters**
- Setting **intercepts** for constrained parameters → fixed via priors





## Specifying mixture models in brms

### Understanding & Identifying parameters of mixture families

#### Step 3: setting priors to identify distributions

```
priors_ZL <-  
  prior(constant(0), class = Intercept, dpar = „mu1“) +  
  prior(constant(0), class = Intercept, dpar = „mu2“) +  
  prior(constant(log(0.0001)), class = Intercept, dpar = „kappa2“)
```

#### Constraints in the Two-Parameter Mixture Model:

1. memory distribution is centered around zero  $\rightarrow \mu_1 = 0$
2. center the guessing distribution around zero  $\rightarrow \mu_2 = 0$
3. guessing distribution is flat  $\rightarrow \kappa_2 \approx 0$  (brms uses a log-link function)
4.  $\theta_2$  is internally fixed to zero by brms




## Specifying mixture models in brms

### Fitting & Summarizing results of mixture models

#### Step 4: estimating the with the *brm* function

1. provide the **specified formula**,
2. names of formula variables = **data** variable names
3. Use **mixture family** as data distribution family
4. submit the **defined priors** to identify your mixture distributions

→ **Model estimation takes some time... be patient.**




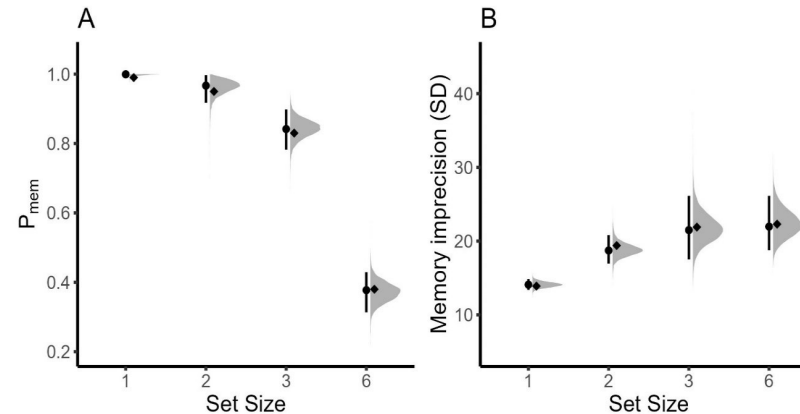
```
fit_ZLmodel <- brm(  
  formula = formula_ZL,  
  data = myData,  
  # now required arguments  
  family = mixFamily_ZL,  
  prior = priors_ZL)
```

# Specifying mixture models in brms

## Fitting & Summarizing results of mixture models

### Step 5: Evaluating model results

1. `pp_check(fit_ZLmodel)` → posterior predictive plot to visually evaluate model fit
2. `summary(fit_ZLmodel)` → overview of the estimated parameters
3. `fixef(fit_ZLmodel)` & `ranef(fit_ZLmodel)` → get fixed and random effect estimates
4. `tidybayes` & `ggplot2` package → processing posterior draws; useful to plot model results



Group-Level Effects:

~subID (Number of levels: 8)

	Estimate	Est.Error	l-95% CI	u-95% CI	Rhat	Bulk_ESS	Tail_ESS
sd(kappa1_setsize1)	0.07	0.06	0.00	0.21	1.00	3844	3490
sd(kappa1_setsize2)	0.21	0.13	0.02	0.51	1.00	2188	2385
sd(kappa1_setsize3)	0.48	0.20	0.20	0.97	1.00	2359	3564
sd(kappa1_setsize6)	0.19	0.16	0.01	0.58	1.00	3848	3370
sd(theta1_setsize1)	1.77	1.54	0.08	5.65	1.00	2326	2845
sd(theta1_setsize2)	1.49	0.70	0.57	3.26	1.00	2908	4425
sd(theta1_setsize3)	0.46	0.27	0.05	1.08	1.00	2633	2557
sd(theta1_setsize6)	0.14	0.12	0.00	0.43	1.00	3549	3506

Population-Level Effects:

	Estimate	Est.Error	l-95% CI	u-95% CI	Rhat	Bulk_ESS	Tail_ESS
mu1_Intercept	0.00	0.00	0.00	0.00	NA	NA	NA
mu2_Intercept	0.00	0.00	0.00	0.00	NA	NA	NA
kappa2_Intercept	-100.00	0.00	-100.00	-100.00	NA	NA	NA
kappa1_setsize1	2.81	0.05	2.70	2.91	1.00	6449	6160
kappa1_setsize2	2.23	0.10	2.01	2.42	1.00	3597	4064
kappa1_setsize3	1.94	0.20	1.55	2.34	1.00	2156	2387
kappa1_setsize6	1.88	0.17	1.55	2.21	1.00	6705	6526
theta1_setsize1	6.46	1.61	4.39	10.78	1.00	2411	1564
theta1_setsize2	3.37	0.68	2.21	4.92	1.00	2532	4005
theta1_setsize3	1.69	0.22	1.26	2.14	1.00	3679	4257
theta1_setsize6	-0.53	0.13	-0.78	-0.28	1.00	5835	4647

# Specifying mixture models in brms

## Fitting & Summarizing results of mixture models

### Step 5: Evaluating model results

**Important:** take care of link function and transformations when interpreting model parameters



→ for computational efficiency and ideal sampling brms transforms parameters with bounded parameter spaces (e.g., precision > 0, 0 > probabilities < 1)

Kappa → log link

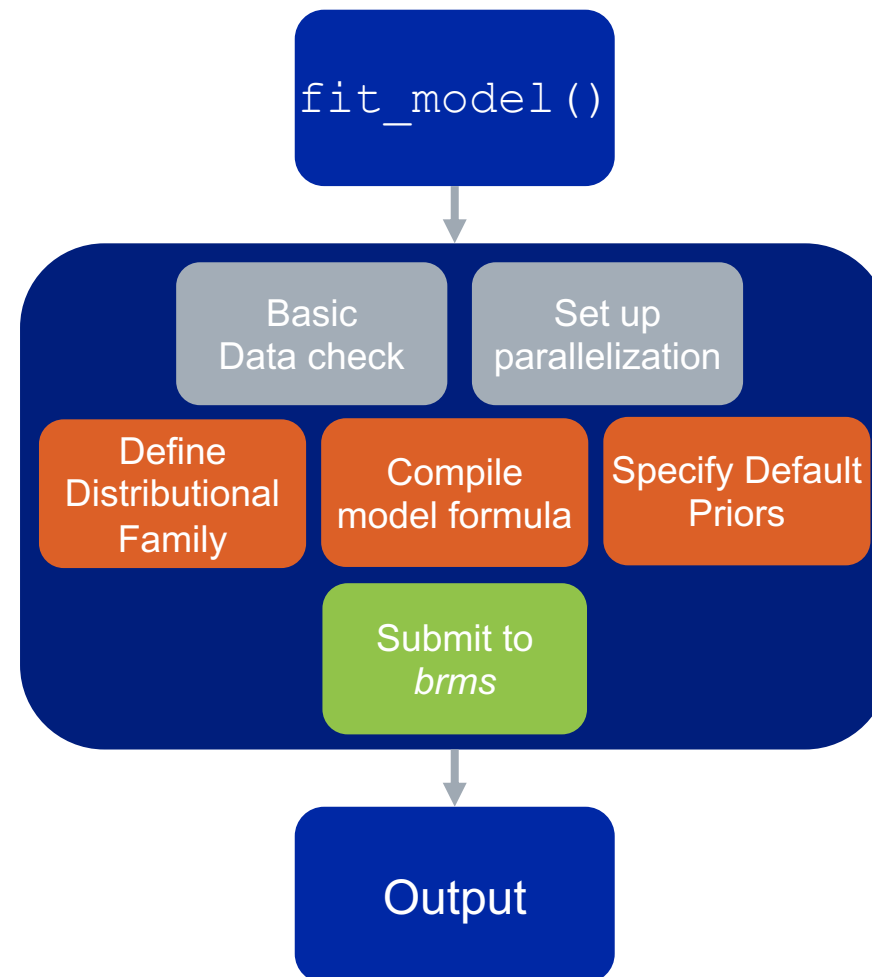
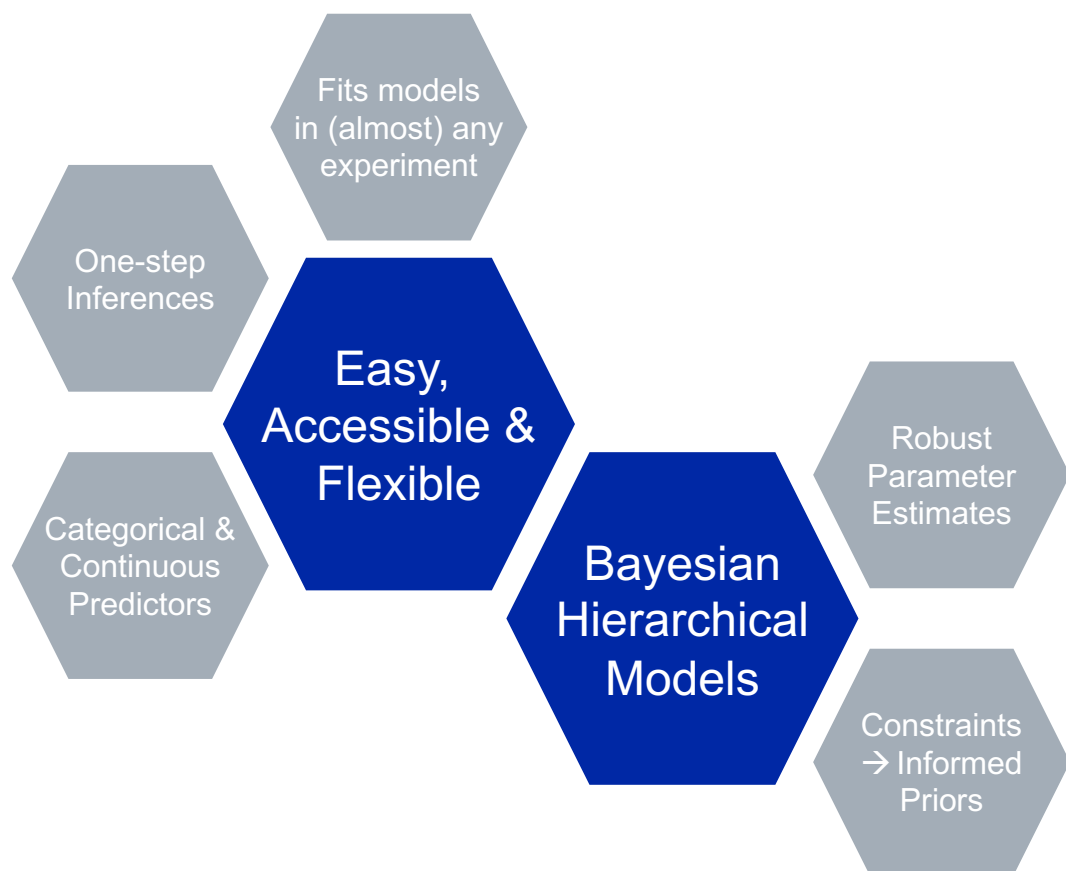
$$\kappa_{native} = e^{\kappa_{par}}; \kappa_{par} = \log(\kappa_{native})$$

Theta → Softmax (for two mixtures = logit)

$$p_i = \frac{e^{\theta_i}}{\sum_{j=1}^K e^{\theta_j}}; p_{Mem} = \frac{e^{\theta_{mem}}}{e^0 + e^{\theta_{mem}}} = \frac{e^{\theta_{mem}}}{1 + e^{\theta_{mem}}}$$



## *bmm*: Easy implementations of mixture models for VWM tasks



## *bmm*: Easy implementations of mixture models for VWM tasks

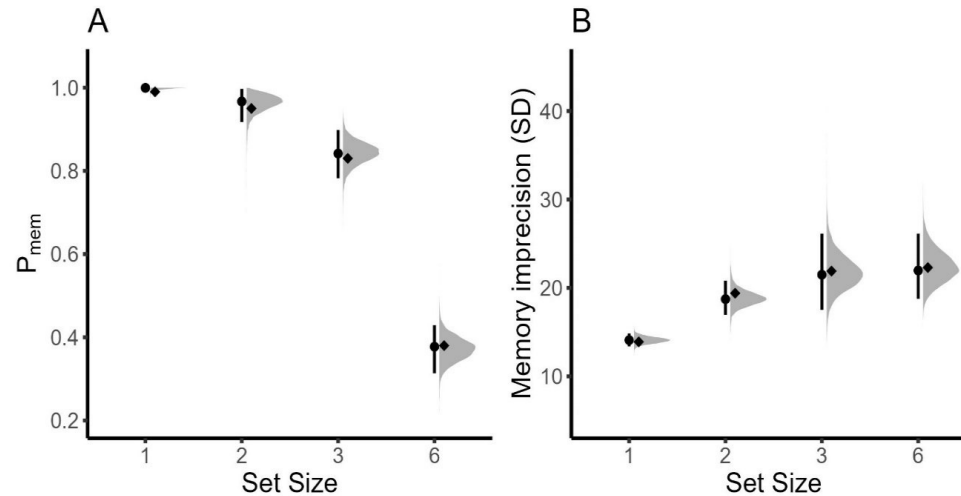
### Two-Parameter Mixture Model

1) Specify  
formula

Predictors  
of Parameters

Dependent  
Variable

```
ZL_mixFormula_bmm <- bf(RespErr ~ 1,  
  kappa ~ 0 + setsize + (0 + setsize || subID),  
  thetat ~ 0 + setsize + (0 + setsize || subID))
```



```
fit_ZL2009_bmm <- bmm::fit_model(formula = ZL_mixFormula,  
  data = data_ZL2008,  
  model type = '2p',  
  warmup=1000, iter=2000, parallel=TRUE)
```

Formula &  
Data

2) Fit Model

Select  
Model Type

Additional  
Arguments

## *bmm*: Easy implementations of mixture models for VWM tasks

### Three-Parameter Mixture Model

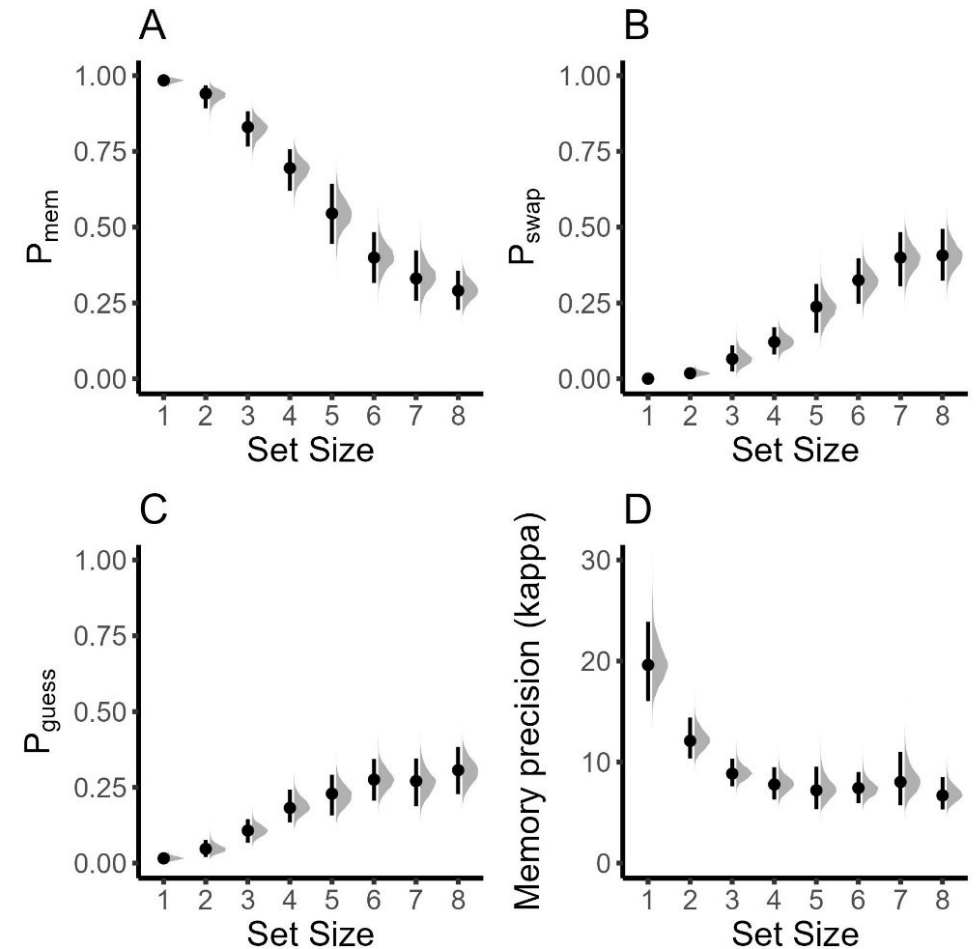
1) Specify  
formula

```
ff <- bf(devRad ~ 1,  
  kappa ~ 0 + SetSize + (0 + SetSize || ID),  
  thetat ~ 0 + SetSize + (0 + SetSize || ID),  
  thetant ~ 0 + SetSize + (0 + SetSize || ID))
```

2) Fit Model

Additional  
Arguments

```
fit_3pMM <- bmm::fit_model(  
  formula = ff,  
  data = df_OberauerLin2017_E1,  
  model type = '3p',  
  non_targets = paste0('Item', 2:8, '_Col_rad'),  
  setsize = "SetSize")
```



# The Bayesian Measurement Model (*bmm*) packages

## Interference Measurement Model

1) Specify  
formula

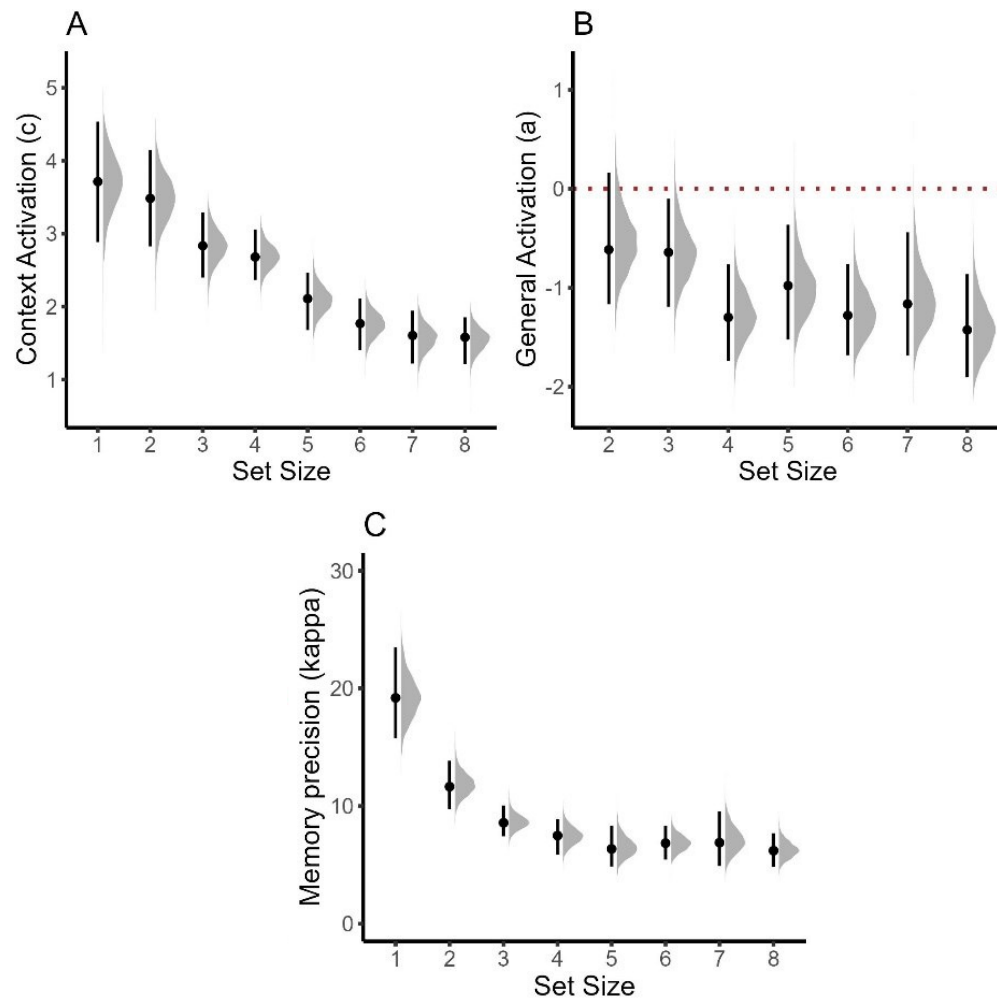
```
ff <- bf(devRad ~ 1,
```

```
  kappa ~ 0 + SetSize + (0 + SetSize || ID),  
  c ~ 0 + SetSize + (0 + SetSize || ID),  
  a ~ 0 + SetSize + (0 + SetSize || ID),  
  s ~ 0 + SetSize + (0 + SetSize || ID),
```

2) Fit Model

Additional  
Arguments

```
fit_IMMfull_mixMod <- fit_model(  
  formula = ff,  
  data = df_OberauerLin2017_E1,  
  model type = 'IMMfull',  
  non_targets = paste0('Item',2:8,'_Col_rad'),  
  spaPos = paste0('Item',2:8,'_Pos_rad'),  
  setsize = "SetSize")
```





# Testing Hypothesis with Bayesian models

## 1. Evaluating 95% Highest Density Intervals

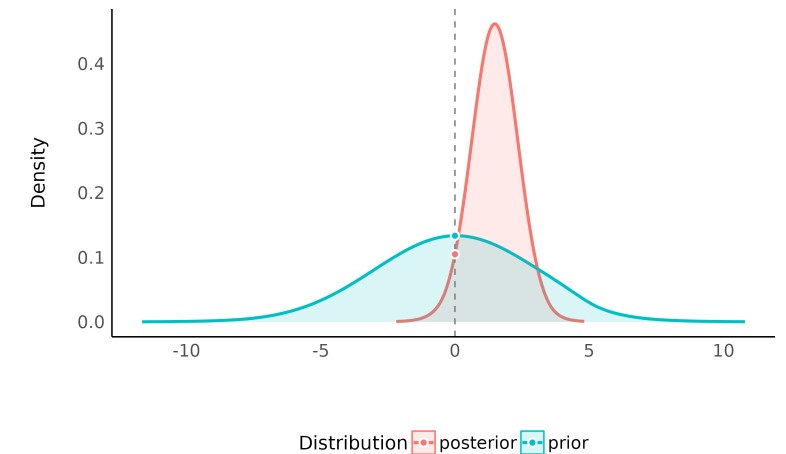

→ Do my posterior estimates, include a specific value (e.g., 0) in their posterior distribution

## 2. Computing Bayes Factors for parameters (Savage-Dickey Ratio)

→ Given the data, has my prior belief credibly changed

## 3. Computing Bayes Factors for competing models (bridgesampling)

→ Under which model are the observed data more probable?







# Time for a coffee break!

Preprint introducing the *bmm* package:

Frischkorn, G. T., & Popov, V. (2023). *A tutorial for estimating mixture models for visual working memory tasks in brms: Introducing the Bayesian Measurement Modeling (bmm) package for R*. PsyArXiv.

<https://doi.org/10.31234/osf.io/umt57>





## Work with (your own) data

### You have your own data

1. Prepare your data
  - a) Transform to long format → each trial in a row
  - b) Calculate response error in radians
  - c) Calculate non-target locations relative to the target values
2. Specify a model you want to fit
3. Test if the modelling is sampling
  - Use a low number of samples to avoid lengthy wait times (iter = 500)
4. See if you can extract and plot results

### You have no own data

1. Choose one of the data sets shared in the GitHub repository (For continuous reproduction tasks, simulated binomial data, etc.)
2. Try to understand the different variables
3. Specify a model you want to fit
4. Test if the model is sampling
  - Use a low number of samples to avoid lengthy wait times (iter = 500)
5. Extract and plot results

## Specifying custom mixture models for accuracy and reaction time data

A logic similar to VWM mixture models can be applied to accuracy data

- Lapses of attention → guessing performance for some trials
- perform with a certain level of ability (i.e., proportion correct)

```
mix_binomial <- mixture(binomial, binomial)
# fix probability to 0.50 via priors
priors_binomial <- prior(
  constant(0), class = Intercept, dpar = "mu2"
)
```

